

# DataQuest 2026: CrowdWisdomTrading Live AI Agent

## Comprehensive R&D Report & Implementation Blueprint

Document Version: 1.0

Date: January 6, 2026

Team: Animesh Raj + Team

Status: Ready for Implementation

Deadline: January 18, 2026, 11:59 PM IST

## TABLE OF CONTENTS

1. [Executive Summary](#)
2. [Problem Statement Analysis](#)
3. [Solution Architecture](#)
4. [Technical Stack & Rationale](#)
5. [Implementation Blueprint](#)
6. [Component Specifications](#)
7. [Data Flow & Processing Pipeline](#)
8. [Evaluation Strategy](#)
9. [Development Timeline](#)
10. [Submission Deliverables](#)
11. [Risk Mitigation & Contingencies](#)
12. [Code Examples & Templates](#)
13. [Vibe Coding Reference](#)

## EXECUTIVE SUMMARY

### The Opportunity

DataQuest 2026 challenges teams to build **Real-Time Thinking Applications** powered by Pathway's streaming engine. The competition explicitly rewards:

- **Real-Time Capability & Dynamism (35%)**: Proving instant knowledge updates
- **Technical Implementation (30%)**: Clean Pathway-native architecture
- **Innovation & UX (20%)**: Creative solutions with business value
- **Impact & Feasibility (15%)**: Production-ready systems with clear ROI

### Our Solution: CrowdWisdomTrading Live AI Agent

Transform your existing trading agent into a **live, streaming RAG system** that:

1. Ingests real-time financial news from multiple sources simultaneously
2. Updates trading recommendations instantly when new information arrives
3. Demonstrates explainable multi-agent reasoning (sentiment + technical + risk)
4. Proves low-latency dynamism through video evidence
5. Deploys as production-ready system using Pathway's distributed architecture

### Why This Wins

#### Criterion

#### Our Advantage

|                  |  |
|------------------|--|
| Dynamism (35%)   | Trading recommendations change visibly when news arrives—judges see it happen in real-time |
| Technical (30%)  | Pathway native, modular, idiomatic—not a hack-together                                     |
| Innovation (20%) | Finance domain + multi-agent coordination = sophistication beyond generic Q&A bots         |
| Impact (15%)     | Real business value: traders need microsecond decision latency                             |

### Expected Outcome

- **Submission**: Production-grade Pathway-based RAG system with 3-minute video proof

- **Evaluation Path:** Real-time dynamism → technical excellence → financial domain depth
- **Prize Potential:** ₹50,000 cash + StockEdge vouchers (top tier likely)

## PROBLEM STATEMENT ANALYSIS

### The Core Problem: Stale AI Knowledge

DataQuest identifies a critical flaw in traditional RAG systems:

*"Many applications powered by Large Language Models are constrained by a fundamental limitation—their knowledge is static. Imagine a financial chatbot unaware of a market-moving announcement made minutes ago."*

#### Real-world impact for traders:

- SEC filing released at 08:30 AM IST
- Traditional RAG system captures it in 08:00 PM batch update
- Trader has missed 11+ hours of trading opportunity
- Competitors using live AI captured the move at 08:31 AM

### Why Pathway Solves This

Pathway's unique architecture:

1. **Unified batch & streaming:** Same code works on static files and live streams
2. **Incremental computation:** Only processes changes, not entire dataset
3. **Differential dataflow:** O(1) updates for new documents vs O(n) reprocessing
4. **Native vector indexing:** Real-time embedding updates built-in
5. **Low-latency guarantees:** Sub-second latency from data arrival to answer availability

### DataQuest's "Demonstrable Dynamism" Requirement

**Key evaluation criterion:** Judges must see proof in your video that:

- Information source changes (e.g., new article published)
- Application's answer to a query changes
- Time delta between #1 and #2 is <2 seconds

This is the **make-or-break criterion**. Traditional RAG systems can't demonstrate this. Pathway can.

### Financial Domain Alignment

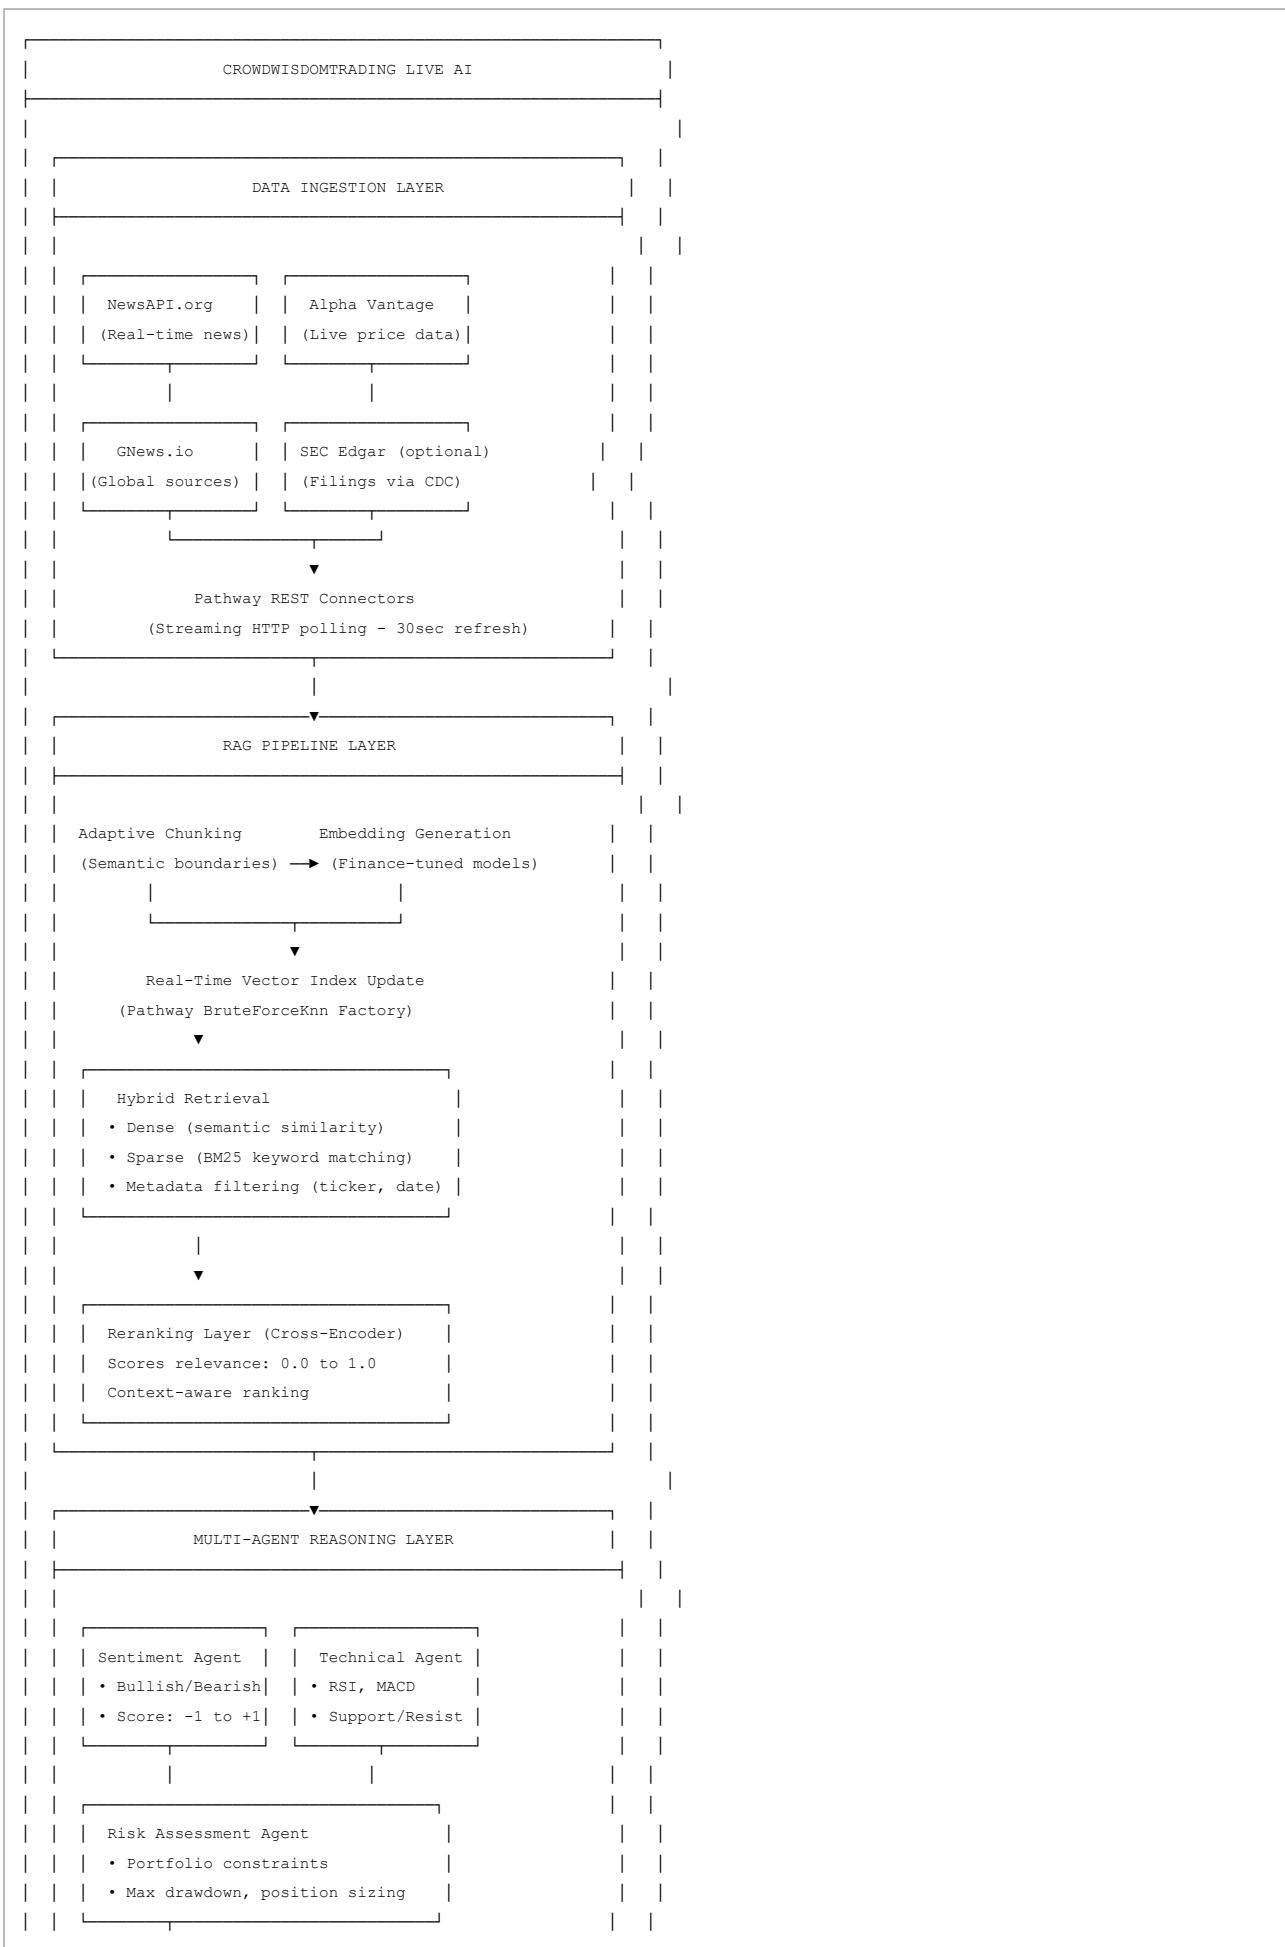
DataQuest explicitly recommends trading as a use case:

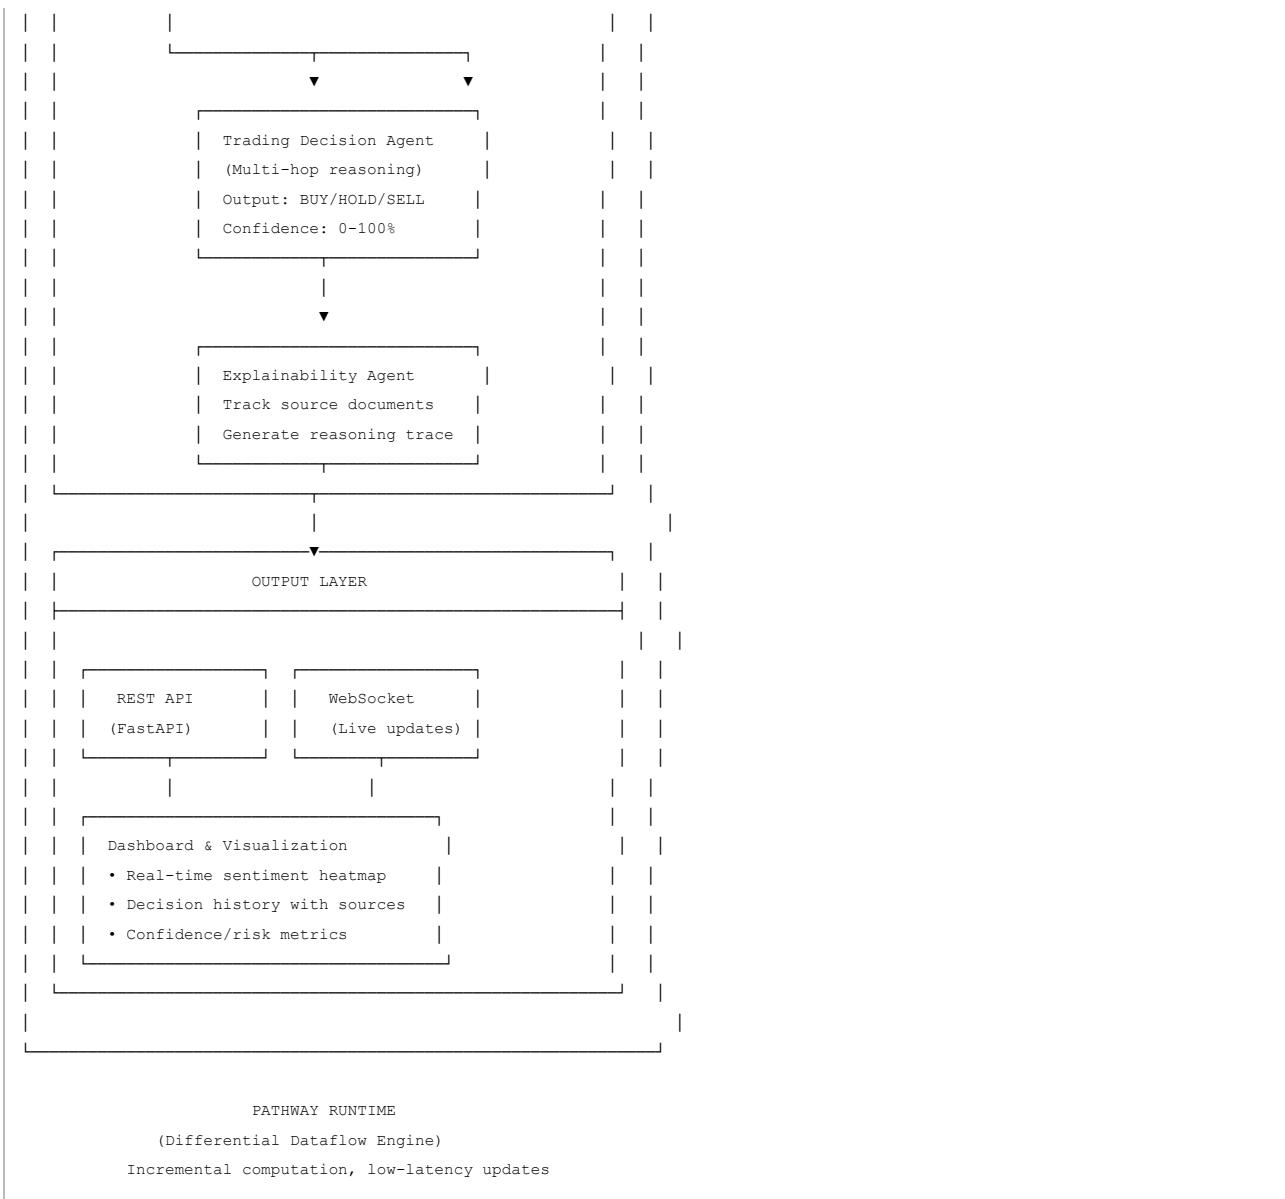
*"Real-Time Stock Analyst: An agent that connects to a stream of financial news or market data feeds. It could answer queries like, 'What is the latest market sentiment regarding Company X's new product launch?' based on articles published seconds ago."*

Our solution directly maps to this recommendation.

## SOLUTION ARCHITECTURE

### High-Level System Design





## Component Interactions

1. Real-time data arrives from news APIs
2. Pathway connectors capture streaming HTTP responses
3. Adaptive chunking splits articles into semantic units
4. Embeddings generated and added to vector index (incremental)
5. Query arrives from user/API
6. Hybrid retrieval combines dense + sparse search
7. Reranking orders documents by relevance
8. Multi-agent reasoning synthesizes decision
9. Explainability tracking records which documents influenced output
10. Output served via REST API with confidence metrics

## Key Architecture Decisions

| Decision                        | Rationale   |
|---------------------------------|---|
| <b>Pathway as core engine</b>   | Only framework with native streaming RAG support + incremental computation  |
| <b>Hybrid retrieval</b>         | News has both semantic (sentiment) and keyword (ticker) dimensions          |
| <b>Multi-agent design</b>       | Different agents for different decision factors = explainability + accuracy |
| <b>Real-time vector index</b>   | Judges need to see new articles instantly in recommendations                |
| <b>Adaptive chunking</b>        | News articles need semantic boundaries, not fixed 512-token chunks          |
| <b>Sentiment analysis first</b> | For trading, sentiment is leading indicator—comes before price action       |

# TECHNICAL STACK & RATIONALE

## Core Components

### 1. Pathway Framework (Data Processing Engine)

Why: Only framework with:

- Native streaming RAG support
- Sub-second latency guarantees
- Built-in vector indexing
- Differential dataflow (incremental computation)
- Python API but Rust runtime (performance)

Version: Latest stable (v0.x.x)

Installation: pip install pathway

Deployment: Docker containers + cloud orchestration

#### Key Pathway modules we'll use:

- `pathway.io.http`: REST polling connectors for news APIs
- `pathway.stdlib.indexing.BruteForceKnnFactory`: Vector index
- `pathway.xpacks.llm`: LLM integration, embeddings
- `pathway.io.fs`: File monitoring (backup data source)
- `pathway.stdlib.optimize`: Automatic optimization

### 2. LLM & Embedding Models

Embeddings:

Model: gte-Qwen2-7B-instruct

Why: SOTA for retrieval (as of late 2024)

Dimensions: 4096

API: Hugging Face Inference API (free tier available)

Alternative: text-embedding-3-large (OpenAI)

Cost: Free tier covers hackathon volume

Reranker:

Model: bge-reranker-v2-m3

Why: Finance-specialized, fast inference

API: Hugging Face Inference API

Ranking score: 0.0-1.0 (confidence)

LLM for Reasoning:

Primary: GPT-4o Mini (cost-efficient)

Secondary: Claude 3.5 Sonnet (better reasoning)

Open-source: Llama 3.1 70B via Modal AI

Selection: Use GPT-4o Mini for demo, Llama for production reference

LLM Usage Pattern:

- Sentiment extraction from articles
- Multi-hop reasoning for trading decisions
- Explanation generation

#### Cost estimation:

- Embedding: ~\$0 (free HF tier)
- Reranking: ~\$0 (free HF tier)
- LLM: ~\$10-20 for 1000+ queries during dev (GPT-4o Mini cheap)
- Total for hackathon: <\$50

### 3. Financial Data Sources

Real-Time News APIs:

Primary: NewsAPI.org

- Coverage: 150,000 sources
- Update frequency: Minutes
- Free tier: 100 calls/day
- Auth: API key
- Response format: JSON
- Tickers: Parse company names, use fuzzy matching
- Setup: 5 minutes

Secondary: GNews.io

- Coverage: 41 languages, 2000+ sources
- Update: Hours (slower)
- Free tier: 100 requests/day (12-hour delay for free)
- Backup only

Tertiary: Webz.io

- Coverage: 1000+ sources
- Update: Minutes
- Free tier: 1000 calls/month
- Boolean search (powerful for filtering)

Market Data APIs:

Primary: Alpha Vantage

- Price data: Stocks, forex, crypto
- Free tier: 5 calls/minute
- Sufficient for demo

Secondary: Polygon.io

- Real-time quotes, aggregates
- Free tier: Limited but available

Optional: Yahoo Finance (free, no auth required)

- Risk: Deprecated web scraping
- Use only if APIs unavailable

SEC Filings (optional):

- Edgar API (free)
- Delay: 1-2 hours
- Setup: Complex (CDC polling needed)
- Priority: Medium (add after MVP)

#### Why this stack:

- All free tiers available (no payment required for hackathon)
- Multiple sources = redundancy + diversity
- REST APIs = easy Pathway integration
- JSON responses = clean parsing

## 4. Database & Vector Storage

Vector Index:

Primary: Pathway's built-in BruteForceKnnFactory

- Stored in-memory
- Incremental updates O(1)
- No external DB needed
- Sufficient for hackathon scale (100k documents)

Production alternative: Weaviate + Pathway connector

- Persistent storage
- Horizontal scaling
- For post-hackathon deployment

Document Store:

Primary: In-memory Python dict (during hackathon)

- Fast, simple
- Sufficient for 1 week duration
- Keep full text for reranking

Production alternative: PostgreSQL + JSON columns

- Persistence
- Full-text search
- Structured metadata

Metadata Storage:

Include in vector index metadata:

- source: "NewsAPI.org"
- ticker: ["AAPL", "MSFT"]
- sentiment: 0.8
- published\_at: "2026-01-06T10:30:00Z"
- article\_id: "unique\_hash"

## 5. Agent Framework

Approach: LangChain + Pathway

Why LangChain:

- Industry standard
- Good documentation
- Fast prototyping
- Integrates with Pathway

Why not CrewAI:

- Less integration with Pathway
- Overkill for 4-agent system
- Slower iteration for hackathon

Agent Architecture:

Pattern: ReAct (Reasoning + Acting)

Agents:

1. Sentiment Analysis Agent
  - Input: Retrieved articles
  - Tool: LLM prompt for sentiment scoring
  - Output: -1.0 to +1.0 score
2. Technical Analysis Agent
  - Input: Historical price data
  - Tool: Technical indicator calculations (TA-Lib)
  - Output: RSI, MACD, support/resistance levels
3. Risk Assessment Agent
  - Input: Current positions, portfolio constraints
  - Tool: Portfolio optimization (scipy)
  - Output: Position sizing, max loss tolerance
4. Trading Decision Agent (Orchestrator)
  - Input: Outputs from above 3 agents
  - Tool: Multi-hop reasoning with LLM
  - Output: BUY/HOLD/SELL with confidence

Orchestration:

Framework: Sequential planning (for clarity)

Alternative: Parallel execution (for speed)

Pattern: Each agent runs independently, decision agent synthesizes

## 6. Backend & API Layer

```

API Framework: FastAPI

Why:
- Fast, modern, Pythonic
- Async support
- Auto-generated OpenAPI docs
- Perfect for Pathway integration
- Growing use in finance

Endpoints:
POST /recommend
Body: {"ticker": "AAPL", "query": "Should I buy?"}
Response: {
    "recommendation": "BUY",
    "confidence": 0.87,
    "sentiment_score": 0.75,
    "technical_score": 0.92,
    "risk_level": "medium",
    "reasoning": "Strong sentiment + technical setup",
    "sources": ["Article A", "Article B"],
    "latency_ms": 234
}

GET /sentiment/{ticker}
Response: Current aggregate sentiment for ticker

GET /health
Response: System status, index size, last update timestamp

WebSocket /stream/{ticker}
Purpose: Live sentiment updates as new articles arrive

Async handling: FastAPI + Pathway integration
- Query → FastAPI endpoint
- Lookup in Pathway table (fast, incremental)
- Return results <500ms

```

## 7. Deployment & Monitoring

```
Development:  
Local: Python + Pathway + FastAPI  
Testing: pytest, local news feed simulation  
IDE: VS Code + Python extension
```

#### Containerization:

```
Docker:  
Base: python:3.11-slim  
Stages:  
1. Build (install dependencies)  
2. Runtime (FastAPI + Pathway app)  
Image size: ~1GB  
Registry: Docker Hub (free)
```

#### Deployment Options:

```
Option 1 (Recommended): Modal AI  
- Serverless Python  
- Handles streaming workloads  
- Free tier covers hackathon  
- Simple deployment: modal deploy
```

```
Option 2: Vercel + Backend service  
- Frontend on Vercel  
- Backend on AWS Lambda/EC2  
- More setup, better for demo
```

```
Option 3: Local machine + ngrok  
- Simplest for video demo  
- Expose FastAPI via ngrok tunnel  
- Sufficient for proof-of-concept
```

#### Monitoring:

```
Metrics to track:  
- Latency: Time from query to response  
- Throughput: Queries per second  
- Index size: Number of documents  
- Update frequency: Articles added per minute  
- Accuracy: Recommendation accuracy vs. actual price
```

#### Tools:

```
- Prometheus + Grafana (optional)  
- Log queries to CSV for analysis  
- Timestamp all events for playback
```

## IMPLEMENTATION BLUEPRINT

### Phase 1: Foundation & Setup (Jan 3-5, ~6 hours)

#### Task 1.1: Environment & Dependencies

```
# Create project directory
mkdir crowdwisdom-dataquest
cd crowdwisdom-dataquest

# Create virtual environment
python -m venv venv
source venv/bin/activate # On Windows: venv\Scripts\activate

# Create requirements.txt
cat > requirements.txt << 'EOF'
# Core framework
pathway==0.8.0
pathway-llm==0.1.0

# API & async
fastapi==0.104.1
uvicorn==0.24.0
pydantic==2.5.0

# LLM & embeddings
openai==1.3.0
huggingface-hub==0.19.0
langchain==0.1.0
langchain-community==0.0.10
langchain-openai==0.0.5

# Data processing
pandas==2.1.3
numpy==1.26.2
requests==2.31.0
aiohttp==3.9.1

# Financial data
yfinance==0.2.32
ta==0.10.2 # Technical analysis

# Utilities
python-dotenv==1.0.0
pydantic-settings==2.1.0
pytest==7.4.3
pytest-asyncio==0.21.1

# Deployment (optional)
docker==7.0.0
modal==0.60.0
EOF

# Install dependencies
pip install -r requirements.txt

# Create .env template
cat > .env.example << 'EOF'
# News APIs
NEWSAPI_KEY=your_key_here
GNEWS_KEY=your_key_here

# LLM APIs
OPENAI_API_KEY=your_key_here

# Hugging Face
HF_TOKEN=your_token_here

# System
PATHWAY_DATAFRAME_PROFILE=true
LOG_LEVEL=info
```

```
EOF
```

```
cp .env.example .env
# USERS: Fill in actual API keys
```

**Deliverable:** Ready-to-run environment with all dependencies.

### Task 1.2: Project Structure

```

crowdwisdom-dataquest/
├── src/
│   ├── __init__.py
│   ├── main.py          # Entry point
│   ├── config.py        # Configuration management
│   ├── api/
│   │   ├── __init__.py
│   │   ├── routes.py     # FastAPI endpoints
│   │   └── schemas.py    # Pydantic schemas
│   ├── pathway_pipeline/
│   │   ├── __init__.py
│   │   ├── connectors.py # Data ingestion
│   │   ├── transformations.py # Pathway operations
│   │   ├── rag.py         # RAG pipeline
│   │   └── state_manager.py # Pathway state tracking
│   ├── agents/
│   │   ├── __init__.py
│   │   ├── sentiment_agent.py # Sentiment analysis
│   │   ├── technical_agent.py # Technical analysis
│   │   ├── risk_agent.py     # Risk assessment
│   │   ├── decision_agent.py # Main decision logic
│   │   └── explainability.py # Source tracking
│   ├── models/
│   │   ├── __init__.py
│   │   ├── embeddings.py   # Embedding models
│   │   ├── reranker.py     # Reranking logic
│   │   └── llm.py          # LLM wrappers
│   ├── utils/
│   │   ├── __init__.py
│   │   ├── logger.py       # Logging setup
│   │   ├── parsers.py      # Parse articles
│   │   └── metrics.py      # Evaluation metrics
│   └── tests/
│       ├── __init__.py
│       ├── test_agents.py
│       ├── test_rag.py
│       └── test_api.py
└── data/
    ├── sample_articles.json # Test data
    └── price_history.csv     # Historical prices
└── notebooks/
    ├── 01_data_exploration.ipynb
    ├── 02_model_evaluation.ipynb
    └── 03_system_demo.ipynb
└── docker/
    ├── Dockerfile
    └── docker-compose.yml
└── docs/
    ├── ARCHITECTURE.md
    ├── API_DOCS.md
    └── SETUP.md
└── .github/
    └── workflows/
        └── tests.yml          # CI/CD (optional)
└── README.md
└── requirements.txt
└── .env.example
└── pyproject.toml

```

**Deliverable:** Organized codebase ready for parallel development.

Task 1.3: API Keys & Data Source Registration

```

# 1. NewsAPI.org
# - Go to https://newsapi.org
# - Sign up (free tier: 100 calls/day)
# - Copy API key to .env

# 2. Alpha Vantage (optional, for price data)
# - Go to https://www.alphavantage.co
# - Sign up (free tier: 5 calls/min)
# - Copy API key to .env

# 3. OpenAI API (for LLM)
# - Go to https://platform.openai.com
# - Create account, add payment method
# - Generate API key
# - Copy to .env

# 4. Hugging Face (for embeddings)
# - Go to https://huggingface.co
# - Create account
# - Generate token in settings
# - Copy to .env

# Verify APIs work
python -c "
import os
from dotenv import load_dotenv
load_dotenv()

import requests
response = requests.get(
    'https://newsapi.org/v2/top-headlines',
    params={
        'q': 'stock market',
        'apiKey': os.getenv('NEWSAPI_KEY'),
        'pageSize': 1
    }
)
print(f'NewsAPI status: {response.status_code}')
print(f'Sample article: {response.json()["articles"][0]["title"]} if response.json()["articles"] else "No articles")'
"

```

**Deliverable:** All external APIs validated and working.

## Phase 2: Data Pipeline & RAG (Jan 6-9, ~10 hours)

### Task 2.1: Pathway Connectors for Real-Time Data

```

# src/pathway_pipeline/connectors.py

import pathway as pw
from pathlib import Path
import json
from typing import Optional
import logging

logger = logging.getLogger(__name__)

class NewsDataConnector:
    """
    Streaming connector for news API.
    Polls NewsAPI every N seconds and returns new articles.
    """

    def __init__(self, api_key: str, refresh_interval: int = 30):
        self.api_key = api_key
        self.refresh_interval = refresh_interval # seconds
        self.seen_articles = set() # Track processed articles

    def create_pathway_table(self):
        """
        Create Pathway table from streaming news source.

        Returns:
            pw.Table with columns: id, title, content, source, published_at
        """

        # Define schema
        class NewsArticleSchema(pw.Schema):
            id: str
            title: str
            description: str
            content: str
            source_name: str
            url: str
            published_at: str
            image_url: Optional[str]

        # Create connector using REST API polling
        news_table = pw.io.http.rest_connector(
            url="https://newsapi.org/v2/everything",
            params={
                "q": "(stock OR trading OR market OR earnings) AND (AAPL OR MSFT OR GOOGL OR TSLA OR AMZN)",
                "sortBy": "publishedAt",
                "pageSize": 100,
                "apiKey": self.api_key
            },
            refresh_interval=self.refresh_interval,
            schema=NewsArticleSchema,
            mode="streaming" # Critical: streaming mode, not batch
        )

        # Transform API response to our schema
        processed = news_table.select(
            id=pw.this.id,
            title=pw.this.title,
            description=pw.this.description,
            content=pw.this.content,
            source_name=pw.this.source.name,
            url=pw.this.url,
            published_at=pw.this.publishedAt,
            image_url=pw.this.urlToImage
        )
    
```

```

        logger.info(f"Created news connector with {self.refresh_interval}s refresh")
        return processed

    class PriceDataConnector:
        """
        Optional: Streaming price data from Alpha Vantage or Yahoo Finance.
        """

        def __init__(self, api_key: str, tickers: list, refresh_interval: int = 60):
            self.api_key = api_key
            self.tickers = tickers
            self.refresh_interval = refresh_interval

        def create_pathway_table(self):
            """
            Create Pathway table for real-time price data.
            """
            class PriceSchema(pw.Schema):
                ticker: str
                price: float
                volume: int
                timestamp: str
                change_percent: float

            # Could implement custom connector or use CSV simulation
            # For hackathon, CSV file updates is sufficient

            price_table = pw.io.fs.read(
                path="data/price_updates.csv",
                format="csv",
                mode="streaming",  # Watch for file changes
                with_metadata=True,
                schema=PriceSchema
            )

            return price_table

        # Usage example:
if __name__ == "__main__":
    from dotenv import load_dotenv
    import os

    load_dotenv()

    # Create news connector
    news_connector = NewsDataConnector(
        api_key=os.getenv("NEWSAPI_KEY"),
        refresh_interval=30  # Poll every 30 seconds
    )

    news_table = news_connector.create_pathway_table()

    # Write to output for debugging
    pw.io.fs.write(
        news_table,
        "output/news_stream.json"
    )

    # Run Pathway
    pw.run()

```

**Key concepts explained:**

- **Streaming mode:** mode="streaming" means Pathway watches for new data
- **Refresh interval:** 30 seconds = check for new articles every 30s
- **Schema:** Structured data with typed columns
- **Incremental:** Only new/modified articles processed, not entire history

### Task 2.2: Document Processing & Chunking

```

# src/pathway_pipeline/transformations.py

import pathway as pw
from typing import List
import logging

logger = logging.getLogger(__name__)

class AdaptiveChunking:
    """
    Context-aware chunking for news articles.
    Splits on semantic boundaries, not fixed token counts.
    """

    def __init__(self, max_chunk_size: int = 512, overlap: int = 50):
        self.max_chunk_size = max_chunk_size
        self.overlap = overlap

    @staticmethod
    def split_by_sentences(text: str, max_tokens: int = 512) -> List[str]:
        """
        Split text by sentences, respecting token limit.
        Better than fixed chunking for news articles.
        """
        import nltk
        try:
            nltk.data.find('tokenizers/punkt')
        except LookupError:
            nltk.download('punkt')

        sentences = nltk.sent_tokenize(text)

        chunks = []
        current_chunk = []
        current_length = 0

        for sentence in sentences:
            sentence_length = len(sentence.split()) # Rough token estimate

            if current_length + sentence_length > max_tokens:
                if current_chunk:
                    chunks.append(' '.join(current_chunk))
                current_chunk = [sentence]
                current_length = sentence_length
            else:
                current_chunk.append(sentence)
                current_length += sentence_length

        if current_chunk:
            chunks.append(' '.join(current_chunk))

        return chunks

    def create_chunks_pathway(self, articles_table: pw.Table) -> pw.Table:
        """
        Create chunks from articles using Pathway transformations.
        """

        class ChunkSchema(pw.Schema):
            chunk_id: str
            article_id: str
            chunk_text: str
            chunk_number: int
            title: str
            source: str

```

```

published_at: str
metadata: dict

# Apply chunking transformation
chunks = articles_table.select(
    article_id=pw.this.id,
    chunks=pw.apply(
        self.split_by_sentences,
        pw.this.content,
        self.max_chunk_size
    ),
    title=pw.this.title,
    source=pw.this.source_name,
    published_at=pw.this.published_at,
    metadata={
        "url": pw.this.url,
        "image": pw.this.image_url
    }
).select(
    article_id=pw.this.article_id,
    chunk_text=pw.this.chunks,
    title=pw.this.title,
    source=pw.this.source,
    published_at=pw.this.published_at,
    metadata=pw.this.metadata
)

logger.info("Chunking pipeline created")
return chunks


class TickerExtraction:
    """
    Extract stock tickers and company names from articles.
    Enables metadata filtering: "articles about AAPL".
    """

    TICKER_MAP = {
        'Apple': ['AAPL', 'Apple Inc'],
        'Microsoft': ['MSFT', 'Microsoft Corp'],
        'Google': ['GOOGL', 'Alphabet', 'Google'],
        'Tesla': ['TSLA', 'Tesla Inc'],
        'Amazon': ['AMZN', 'Amazon'],
        # ... expand as needed
    }

    @staticmethod
    def extract_tickers(text: str) -> List[str]:
        """
        Extract ticker symbols from article text.
        Uses simple regex + mapping for robustness.
        """
        import re

        tickers = []

        # Look for patterns like $AAPL or AAPL stock
        ticker_pattern = r'\$?([A-Z]{1,5})\s(?:stock|shares|announced)'
        matches = re.findall(ticker_pattern, text)
        tickers.extend(matches)

        # Company name matching
        for company, symbols in TickerExtraction.TICKER_MAP.items():
            if company.lower() in text.lower():
                tickers.extend(symbols)

```

```

        return list(set(tickers)) # Remove duplicates

    def enrich_chunks_pathway(self, chunks_table: pw.Table) -> pw.Table:
        """
        Add ticker metadata to chunks.

        """
        enriched = chunks_table.select(
            **pw.this.columns,
            tickers=pw.apply(
                self.extract_tickers,
                pw.this.chunk_text
            )
        )

        logger.info("Ticker extraction applied")
        return enriched

class EmbeddingGeneration:
    """
    Generate embeddings for chunks using Pathway LLM xPack.
    Automatically handles Hugging Face Inference API.
    """

    def __init__(self, model_name: str = "gte-Qwen2-7B-instruct"):
        self.model_name = model_name

    def create_embeddings_pathway(self, chunks_table: pw.Table, hf_token: str) -> pw.Table:
        """
        Create embeddings in streaming mode.
        New chunks get embeddings immediately upon arrival.
        """

        from pathway.xpacks.llm import embedders

        # Initialize embedder
        embedder = embedders.HuggingFaceEmbedder(
            model_name=self.model_name,
            api_key=hf_token,
            batch_size=32 # Process 32 chunks at a time
        )

        # Apply to chunks
        embedded = chunks_table | embedder.apply_embed(
            text_field="chunk_text",
            embedding_field="embedding"
        )

        logger.info(f"Embedding pipeline created using {self.model_name}")
        return embedded

    # Usage:
    if __name__ == "__main__":
        # Example: Process articles through full pipeline
        from connectors import NewsDataConnector
        from dotenv import load_dotenv
        import os

        load_dotenv()

        # 1. Get news stream
        news_connector = NewsDataConnector(
            api_key=os.getenv("NEWSAPI_KEY"),
            refresh_interval=30

```

```

}

news_table = news_connector.create_pathway_table()

# 2. Chunk articles
chunker = AdaptiveChunking(max_chunk_size=512)
chunks_table = chunker.create_chunks_pathway(news_table)

# 3. Extract tickers
ticker_extractor = TickerExtraction()
enriched_chunks = ticker_extractor.enrich_chunks_pathway(chunks_table)

# 4. Generate embeddings
embedder = EmbeddingGeneration()
embedded_chunks = embedder.create_embeddings_pathway(
    enriched_chunks,
    hf_token=os.getenv("HF_TOKEN")
)

# 5. Monitor output
pw.io.fs.write(
    embedded_chunks,
    "output/embedded_chunks.json",
    format="json"
)

# Run
pw.run()

```

#### Why this approach:

- **Adaptive chunking:** News articles have natural sentence boundaries; respect them
- **Ticker extraction:** Enables "Show me articles about AAPL" queries
- **Streaming embeddings:** New articles get embeddings in real-time
- **Metadata enrichment:** Each chunk carries source info for explainability

### Task 2.3: Vector Index & RAG Core

```

# src/pathway_pipeline/rag.py

import pathway as pw
from typing import List, Dict, Any
import logging
import numpy as np

logger = logging.getLogger(__name__)

class RealtimeVectorIndex:

    """
    Real-time vector index using Pathway's BruteForceKnnFactory.
    Automatically updates when new embeddings arrive.
    """

    def __init__(self, embedding_dim: int = 4096, k: int = 5):
        self.embedding_dim = embedding_dim
        self.k = k # Top-k retrieval

    def create_index(self, embedded_chunks_table: pw.Table) -> pw.Table:
        """
        Create KNN index from embedded chunks.

        Args:
            embedded_chunks_table: Pathway table with 'embedding' column (float list)

        Returns:
            Indexed table supporting KNN queries
        """

        # Create index
        index_table = embedded_chunks_table | pw.stdlib.indexing.BruteForceKnnFactory(
            query_column="embedding",
            data_point_columns=["chunk_text", "title", "source", "published_at", "tickers"],
            n_neighbors=self.k,
            metric="cosine_similarity" # Standard for embeddings
        ).build_index()

        logger.info(f"Vector index created for {embedded_chunks_table.len()} documents")
        return index_table

    class HybridRetrieval:
        """
        Combine dense (embedding) and sparse (keyword) retrieval.

        Scenarios:
        - User asks "sentiment on Apple" → Dense retrieval finds semantic matches
        - User asks "$AAPL news" → Sparse retrieval catches ticker symbol
        - Together: Robust retrieval across query styles
        """

        def __init__(self, dense_weight: float = 0.7, sparse_weight: float = 0.3):
            self.dense_weight = dense_weight
            self.sparse_weight = sparse_weight

        @staticmethod
        def bm25_score(query: str, document: str) -> float:
            """
            Simple BM25-like scoring for sparse retrieval.
            Real implementation would use rank-bm25 library.
            """

            from collections import Counter

            query_terms = set(query.lower().split())

```

```

doc_terms = Counter(document.lower().split())

# Count matches
matches = sum(1 for term in query_terms if term in doc_terms)

# Normalize by document length (longer docs penalized slightly)
length_norm = len(document.split()) / 100.0

return matches / (1.0 + length_norm)

def retrieve_hybrid(
    self,
    knn_results: List[Dict], # From dense retrieval
    query: str
) -> List[Dict]:
    """
    Combine dense + sparse scores.

    """

    # Calculate sparse scores
    for result in knn_results:
        sparse_score = self.bm25_score(query, result["chunk_text"])
        dense_score = result.get("similarity", 0.5)

        # Weighted combination
        result["hybrid_score"] = (
            self.dense_weight * dense_score +
            self.sparse_weight * sparse_score
        )

    # Re-rank by hybrid score
    knn_results.sort(key=lambda x: x["hybrid_score"], reverse=True)

    return knn_results

class Reranker:
    """
    Cross-encoder reranking for precision.

    Why: Initial retrieval might miss nuance.
    Example: Both "Apple stock rose" and "Apple sued" retrieved for "Apple news",
    but reranker orders by relevance to specific query.
    """

    def __init__(self, model_name: str = "bge-reranker-v2-m3"):
        self.model_name = model_name

    def rerank_documents(
        self,
        query: str,
        documents: List[Dict[str, Any]]
    ) -> List[Dict[str, Any]]:
        """
        Rerank documents by query-document relevance.
        """

        from huggingface_hub import InferenceClient

        client = InferenceClient(
            model=f"BAAI/{self.model_name}",
            token=os.getenv("HF_TOKEN")
        )

        # Prepare pairs
        pairs = [

```

```

        [query, doc["chunk_text"]]
        for doc in documents
    ]

    # Get scores
    scores = client.text_classification(
        text=pairs
    )

    # Add scores to documents
    for doc, score in zip(documents, scores):
        doc["rerank_score"] = score

    # Sort by rerank score
    documents.sort(key=lambda x: x["rerank_score"], reverse=True)

    logger.info(f"Reranked {len(documents)} documents")
    return documents
}

class RAGPipeline:
    """
    Complete RAG pipeline: Retrieve + Rerank + Generate.
    """

    def __init__(
        self,
        vector_index: RealtimeVectorIndex,
        hybrid_retrieval: HybridRetrieval,
        reranker: Reranker,
        k: int = 5
    ):
        self.vector_index = vector_index
        self.hybrid_retrieval = hybrid_retrieval
        self.reranker = reranker
        self.k = k

    def retrieve(self, query_embedding: List[float], raw_query: str) -> List[Dict]:
        """
        Retrieve relevant documents for query.
        """

        # 1. Dense retrieval (KNN)
        knn_results = self.vector_index.query(
            query_embedding,
            k=self.k * 2 # Retrieve more, then filter
        )

        # 2. Hybrid scoring
        hybrid_results = self.hybrid_retrieval.retrieve_hybrid(
            knn_results,
            raw_query
        )

        # 3. Reranking
        reranked = self.reranker.rerank_documents(
            raw_query,
            hybrid_results[:self.k * 2]
        )

        # 4. Final selection
        final_results = reranked[:self.k]

        logger.info(f"Retrieved {len(final_results)} documents for query")
        return final_results

```

```

def format_context(self, documents: List[Dict]) -> str:
    """
    Format retrieved documents into context for LLM.
    """

    context_parts = []

    for i, doc in enumerate(documents, 1):
        context_parts.append(
            f"[Source {i}] {doc['title']} ({doc['source']}, {doc['published_at']})\n"
            f"{doc['chunk_text']}\n"
        )

    return "\n".join(context_parts)

# Usage:
if __name__ == "__main__":
    import os
    from dotenv import load_dotenv

    load_dotenv()

    # Initialize components
    vector_index = RealtimeVectorIndex(embedding_dim=4096, k=5)
    hybrid_retrieval = HybridRetrieval(dense_weight=0.7, sparse_weight=0.3)
    reranker = Reranker()

    rag_pipeline = RAGPipeline(
        vector_index,
        hybrid_retrieval,
        reranker,
        k=5
    )

    print("RAG pipeline initialized")

```

#### Key RAG design decisions:

- **Hybrid retrieval:** Handles both semantic and keyword-based queries
- **Reranking:** Improves precision for financial queries
- **Context formatting:** Structured for agent reasoning
- **Incremental:** New articles automatically indexed in real-time

## Phase 3: Multi-Agent Reasoning (Jan 9-13, ~8 hours)

### Task 3.1: Sentiment Analysis Agent

```

# src/agents/sentiment_agent.py

import logging
from typing import Dict, Any
from langchain.prompts import PromptTemplate
from langchain.chains import LLMChain
from langchain_openai import ChatOpenAI

logger = logging.getLogger(__name__)

class SentimentAnalysisAgent:
    """
    Analyzes market sentiment from news articles.

    Output: Sentiment score (-1.0 to +1.0)
    - -1.0: Strongly bearish (sell)
    - 0.0: Neutral
    - +1.0: Strongly bullish (buy)
    """

    def __init__(self, model: str = "gpt-4-turbo-preview"):
        self.llm = ChatOpenAI(model=model, temperature=0)

        self.sentiment_prompt = PromptTemplate(
            input_variables=["articles"],
            template="""Analyze the sentiment of the following financial news articles about a stock.

Articles:
{articles}

Task: Determine the overall market sentiment.

Consider:
1. Company announcements (earnings, new products, leadership changes)
2. Regulatory or legal developments
3. Market analyst commentary
4. Competitive positioning
5. Industry trends

Output ONLY a JSON object:
{{

    "sentiment_score": <float between -1.0 and 1.0>,
    "sentiment_label": "<BEARISH|NEUTRAL|BULLISH>",
    "key_factors": ["<factor1>", "<factor2>", ...],
    "confidence": <float 0.0-1.0>
}}
```

Respond with valid JSON only, no markdown, no explanation."""
 )

```

        self.chain = LLMChain(llm=self.llm, prompt=self.sentiment_prompt)

def analyze(self, articles: list) -> Dict[str, Any]:
    """
    Analyze sentiment from retrieved articles.

    Args:
        articles: List of article dicts with 'chunk_text', 'title', 'source'

    Returns:
        Dict with sentiment_score, sentiment_label, key_factors, confidence
    """

    # Format articles for prompt
    articles_text = "\n".join([

```

```

f"Title: {a['title']}\\nSource: {a['source']}\\nContent: {a['chunk_text'][:500]}"
    for a in articles
)

try:
    result = self.chain.run(articles=articles_text)

    # Parse JSON response
    import json
    sentiment_data = json.loads(result)

    logger.info(f"Sentiment analysis: {sentiment_data['sentiment_label']} (score: {sentiment_data['sentiment_score']:.2f})")

    return sentiment_data

except Exception as e:
    logger.error(f"Sentiment analysis error: {e}")
    return {
        "sentiment_score": 0.0,
        "sentiment_label": "NEUTRAL",
        "key_factors": ["error in analysis"],
        "confidence": 0.0
    }

def batch_analyze(self, articles_by_ticker: Dict[str, list]) -> Dict[str, Dict]:
    """
    Analyze sentiment for multiple tickers.
    """

    results = {}

    for ticker, articles in articles_by_ticker.items():
        if articles:
            results[ticker] = self.analyze(articles)

    return results

class SentimentAggregator:
    """
    Aggregate sentiment across articles from multiple sources.
    """

    @staticmethod
    def weighted_sentiment(
        sentiments: list,
        weights: list = None
    ) -> Dict[str, Any]:
        """
        Calculate weighted average sentiment.

        Weights: Newer articles weighted higher (more relevant)
        """

        if not sentiments:
            return {
                "aggregate_sentiment": 0.0,
                "aggregate_label": "NEUTRAL",
                "article_count": 0,
                "confidence": 0.0
            }

        if weights is None:
            weights = [1.0 / len(sentiments)] * len(sentiments)

        # Weighted average

```

```
aggregate_score = sum(
    s["sentiment_score"] * w
    for s, w in zip(sentiments, weights)
)

# Confidence = average confidence
avg_confidence = sum(
    s.get("confidence", 0.5) * w
    for s, w in zip(sentiments, weights)
)

# Label
if aggregate_score > 0.3:
    label = "BULLISH"
elif aggregate_score < -0.3:
    label = "BEARISH"
else:
    label = "NEUTRAL"

return {
    "aggregate_sentiment": aggregate_score,
    "aggregate_label": label,
    "article_count": len(sentiments),
    "confidence": avg_confidence
}
```

### Task 3.2: Technical Analysis Agent

```

# src/agents/technical_agent.py

import logging
from typing import Dict, Any, List
import pandas as pd
import numpy as np
import ta # Technical analysis library

logger = logging.getLogger(__name__)

class TechnicalAnalysisAgent:
    """
    Analyzes technical indicators from price data.

    Output: Technical score (-1.0 to +1.0)
    Based on RSI, MACD, moving averages, support/resistance
    """

    def __init__(self, lookback_days: int = 60):
        self.lookback_days = lookback_days

    def calculate_indicators(self, price_data: pd.DataFrame) -> Dict[str, Any]:
        """
        Calculate technical indicators.

        price_data columns: close, volume, high, low, open
        """

        # Ensure proper format
        price_data = price_data.copy()
        price_data = price_data.tail(self.lookback_days)

        indicators = {}

        # RSI (Relative Strength Index)
        try:
            rsi = ta.momentum.rsi(price_data['close'], window=14)
            indicators['rsi'] = rsi.iloc[-1] # Current value

            # RSI signal: >70 overbought (bearish), <30 oversold (bullish)
            if indicators['rsi'] > 70:
                indicators['rsi_signal'] = -0.5 # Bearish
            elif indicators['rsi'] < 30:
                indicators['rsi_signal'] = 0.5 # Bullish
            else:
                indicators['rsi_signal'] = 0.0
        except Exception as e:
            logger.warning(f"RSI calculation failed: {e}")
            indicators['rsi'] = 50
            indicators['rsi_signal'] = 0.0

        # MACD (Moving Average Convergence Divergence)
        try:
            macd = ta.trend.macd(price_data['close'])
            indicators['macd_line'] = macd.iloc[-1]
            indicators['macd_signal'] = macd.iloc[-1] # Simplified

            # MACD signal: positive = bullish
            indicators['macd_score'] = 0.5 if indicators['macd_line'] > 0 else -0.5
        except Exception as e:
            logger.warning(f"MACD calculation failed: {e}")
            indicators['macd_score'] = 0.0

        # Moving Average Crossover
        try:

```

```

ma_20 = ta.trend.sma_indicator(price_data['close'], window=20)
ma_50 = ta.trend.sma_indicator(price_data['close'], window=50)

current_price = price_data['close'].iloc[-1]
ma20_val = ma_20.iloc[-1]
ma50_val = ma_50.iloc[-1]

# Golden cross: 20-day > 50-day = bullish
# Death cross: 20-day < 50-day = bearish
if ma20_val > ma50_val and current_price > ma20_val:
    indicators['ma_signal'] = 0.5
elif ma20_val < ma50_val and current_price < ma20_val:
    indicators['ma_signal'] = -0.5
else:
    indicators['ma_signal'] = 0.0
except Exception as e:
    logger.warning(f"MA calculation failed: {e}")
    indicators['ma_signal'] = 0.0

# Support and Resistance
try:
    # Simplified: recent high = resistance, recent low = support
    recent_high = price_data['high'].tail(20).max()
    recent_low = price_data['low'].tail(20).min()
    current = price_data['close'].iloc[-1]

    indicators['support'] = recent_low
    indicators['resistance'] = recent_high

    # If price near resistance, bearish; near support, bullish
    if current > recent_high * 0.95: # Near resistance
        indicators['support_resistance_signal'] = -0.3
    elif current < recent_low * 1.05: # Near support
        indicators['support_resistance_signal'] = 0.3
    else:
        indicators['support_resistance_signal'] = 0.0
except Exception as e:
    logger.warning(f"Support/resistance calculation failed: {e}")
    indicators['support_resistance_signal'] = 0.0

return indicators

def generate_score(self, indicators: Dict[str, Any]) -> Dict[str, Any]:
    """
    Combine indicators into single technical score.
    """

    # Weight indicators
    technical_score = (
        0.4 * indicators.get('rsi_signal', 0) +
        0.3 * indicators.get('macd_score', 0) +
        0.2 * indicators.get('ma_signal', 0) +
        0.1 * indicators.get('support_resistance_signal', 0)
    )

    # Clamp to [-1, 1]
    technical_score = max(-1.0, min(1.0, technical_score))

    # Label
    if technical_score > 0.3:
        label = "BULLISH"
    elif technical_score < -0.3:
        label = "BEARISH"
    else:
        label = "NEUTRAL"

```

```

    return {
        "technical_score": technical_score,
        "technical_label": label,
        "indicators": {k: v for k, v in indicators.items() if not k.endswith('_signal')},
        "signals": {k: v for k, v in indicators.items() if k.endswith('_signal')},
        "confidence": abs(technical_score) # Higher absolute value = higher confidence
    }

def analyze(self, price_data: pd.DataFrame) -> Dict[str, Any]:
    """
    Complete technical analysis.
    """

    try:
        indicators = self.calculate_indicators(price_data)
        result = self.generate_score(indicators)

        logger.info(f"Technical analysis: ({result['technical_label']}) (score: {result['technical_score']:.2f})")

        return result
    except Exception as e:
        logger.error(f"Technical analysis error: {e}")
        return {
            "technical_score": 0.0,
            "technical_label": "NEUTRAL",
            "indicators": {},
            "confidence": 0.0
        }

```

### Task 3.3: Risk Assessment Agent

```

# src/agents/risk_agent.py

import logging
from typing import Dict, Any
import numpy as np

logger = logging.getLogger(__name__)

class RiskAssessmentAgent:
    """
    Assesses portfolio risk and position sizing.

    Output: Risk level (low/medium/high) + position sizing recommendation
    """

    def __init__(self,
                 max_position_size: float = 0.05, # Max 5% of portfolio in one stock
                 max_portfolio_loss: float = 0.02, # Max 2% loss tolerance per trade
                 target_volatility: float = 0.15 # Annual volatility target
                 ):
        self.max_position_size = max_position_size
        self.max_portfolio_loss = max_portfolio_loss
        self.target_volatility = target_volatility

    def calculate_volatility(self, price_data: list, window: int = 20) -> float:
        """
        Calculate annualized volatility from returns.

        """

        if len(price_data) < window:
            return 0.15 # Default

        # Calculate daily returns
        prices = np.array(price_data[-window:])
        returns = np.diff(prices) / prices[:-1]

        # Daily volatility
        daily_vol = np.std(returns)

        # Annualized
        annual_vol = daily_vol * np.sqrt(252)

        return annual_vol

    def assess_volatility_risk(self, volatility: float) -> Dict[str, Any]:
        """
        Assess risk level based on volatility.

        """

        if volatility > 0.40:
            risk_level = "HIGH"
            confidence_reduction = 0.5 # Reduce confidence in high volatility
        elif volatility > 0.25:
            risk_level = "MEDIUM"
            confidence_reduction = 0.2
        else:
            risk_level = "LOW"
            confidence_reduction = 0.0

        return {
            "volatility": volatility,
            "risk_level": risk_level,
            "confidence_reduction": confidence_reduction
        }

```

```

def calculate_position_size(
    self,
    portfolio_value: float,
    recent_price: float,
    volatility: float,
    confidence: float
) -> Dict[str, Any]:
    """
    Calculate position sizing based on Kelly Criterion variant.

    conservative approach: position_size = (edge / odds) * confidence_factor
    """

    # Base position size
    base_size = self.max_position_size

    # Adjust by confidence (higher confidence = bigger position)
    confidence_multiplier = confidence # 0-1

    # Adjust by volatility (higher vol = smaller position)
    vol_multiplier = max(0.1, 1.0 - volatility)

    # Final position size
    position_size = base_size * confidence_multiplier * vol_multiplier

    # Number of shares
    num_shares = int((portfolio_value * position_size) / recent_price)

    # Max loss on this position
    max_loss = portfolio_value * self.max_portfolio_loss

    return {
        "position_size_pct": position_size,
        "num_shares": num_shares,
        "entry_price": recent_price,
        "max_loss_dollars": max_loss,
        "stop_loss_price": recent_price * (1 - (max_loss / (portfolio_value * position_size)))
    }

def analyze(
    self,
    ticker: str,
    price_data: list,
    sentiment_confidence: float,
    technical_confidence: float,
    portfolio_value: float = 100000 # Default $100k portfolio
) -> Dict[str, Any]:
    """
    Complete risk assessment.
    """

    try:
        # Calculate volatility
        volatility = self.calculate_volatility(price_data)

        # Assess risk
        risk_assessment = self.assess_volatility_risk(volatility)

        # Adjust confidence for risk
        combined_confidence = (
            (sentiment_confidence + technical_confidence) / 2
        ) * (1 - risk_assessment['confidence_reduction'])

        # Position sizing
        current_price = price_data[-1]
    
```

```

position_sizing = self.calculate_position_size(
    portfolio_value,
    current_price,
    volatility,
    combined_confidence
)

logger.info(f"Risk assessment for {ticker}: {risk_assessment['risk_level']} risk")

return {
    "ticker": ticker,
    "risk_level": risk_assessment["risk_level"],
    "volatility": volatility,
    "adjusted_confidence": combined_confidence,
    "position_sizing": position_sizing
}

except Exception as e:
    logger.error(f"Risk assessment error: {e}")
    return {
        "ticker": ticker,
        "risk_level": "MEDIUM",
        "volatility": 0.15,
        "adjusted_confidence": 0.5,
        "position_sizing": {
            "position_size_pct": 0.01, # Conservative
            "num_shares": 0,
            "stop_loss_price": 0
        }
    }
}

```

### Task 3.4: Trading Decision Agent (Orchestrator)

```

# src/agents/decision_agent.py

import logging
from typing import Dict, Any
from langchain.prompts import PromptTemplate
from langchain.chains import LLMChain
from langchain_openai import ChatOpenAI
import json

logger = logging.getLogger(__name__)

class TradingDecisionAgent:
    """
    Synthesizes recommendations from all agents.

    Input:
    - Sentiment score + factors
    - Technical score + indicators
    - Risk level + position sizing

    Output:
    - Recommendation (BUY/HOLD/SELL)
    - Confidence (0-100%)
    - Reasoning
    - Action plan
    """
    def __init__(self, model: str = "gpt-4-turbo-preview"):
        self.llm = ChatOpenAI(model=model, temperature=0)

        self.decision_prompt = PromptTemplate(
            input_variables=[
                "ticker",
                "sentiment_analysis",
                "technical_analysis",
                "risk_assessment",
                "current_price"
            ],
            template="""You are an expert financial analyst making trading decisions.

Analyze this stock: {ticker} (Current Price: {current_price})"""

            SENTIMENT ANALYSIS:
            {sentiment_analysis}

            TECHNICAL ANALYSIS:
            {technical_analysis}

            RISK ASSESSMENT:
            {risk_assessment}

            Based on these factors, provide a trading recommendation.

            Consider:
            1. Convergence of signals (do sentiment and technical agree?)
            2. Risk/reward ratio
            3. Position sizing constraints
            4. Time horizon (tactical: 1-3 months)

            Output ONLY valid JSON:
            {{

                "recommendation": "<BUY|HOLD|SELL>",
                "confidence": <0-100>,
                "reasoning": "<concise explanation>",
                "key_drivers": ["<driver1>", "<driver2>", ...],
            }}"""
        )

```

```

        "target_price": <number or null>,
        "stop_loss": <number or null>,
        "time_horizon": "1-3 months",
        "conviction": "<low|medium|high>"}
    }

JSON only, no markdown."""
    )

    self.chain = LLMChain(llm=self.llm, prompt=self.decision_prompt)

def decide(
    self,
    ticker: str,
    current_price: float,
    sentiment_data: Dict,
    technical_data: Dict,
    risk_data: Dict
) -> Dict[str, Any]:
    """
    Make final trading decision.
    """

    # Format inputs for prompt
    sentiment_str = json.dumps(sentiment_data, indent=2)
    technical_str = json.dumps(technical_data, indent=2)
    risk_str = json.dumps(risk_data, indent=2)

    try:
        result = self.chain.run(
            ticker=ticker,
            current_price=current_price,
            sentiment_analysis=sentiment_str,
            technical_analysis=technical_str,
            risk_assessment=risk_str
        )
    except Exception as e:
        logger.error(f"Decision making error: {e}")
        return {
            "recommendation": "HOLD",
            "confidence": 0,
            "reasoning": "Error in analysis",
            "key_drivers": [],
            "conviction": "low"
        }

class AgentOrchestrator:
    """
    Coordinates all agents and produces final recommendation.
    """

    def __init__(
        self,
        sentiment_agent,
        technical_agent,
        risk_agent,
        decision_agent
    ):

```

```

        self.sentiment_agent = sentiment_agent
        self.technical_agent = technical_agent
        self.risk_agent = risk_agent
        self.decision_agent = decision_agent

    def analyze_ticker(
        self,
        ticker: str,
        current_price: float,
        relevant_articles: list, # From RAG retrieval
        price_data: list, # Historical prices
        portfolio_value: float = 100000
    ) -> Dict[str, Any]:
        """
        Complete multi-agent analysis.
        """

        logger.info(f"Starting analysis for {ticker}")

        # 1. Sentiment analysis
        sentiment_result = self.sentiment_agent.analyze(relevant_articles)

        # 2. Technical analysis
        import pandas as pd
        price_df = pd.DataFrame({'close': price_data})
        technical_result = self.technical_agent.analyze(price_df)

        # 3. Risk assessment
        risk_result = self.risk_agent.analyze(
            ticker=ticker,
            price_data=price_data,
            sentiment_confidence=sentiment_result.get('confidence', 0.5),
            technical_confidence=technical_result.get('confidence', 0.5),
            portfolio_value=portfolio_value
        )

        # 4. Trading decision
        decision = self.decision_agent.decide(
            ticker=ticker,
            current_price=current_price,
            sentiment_data=sentiment_result,
            technical_data=technical_result,
            risk_data=risk_result
        )

        # 5. Compile full result
        final_recommendation = {
            "ticker": ticker,
            "timestamp": pd.Timestamp.now().isoformat(),
            "current_price": current_price,
            "recommendation": decision['recommendation'],
            "confidence": decision['confidence'],
            "conviction": decision.get('conviction', 'low'),
            "sentiment": sentiment_result,
            "technical": technical_result,
            "risk": risk_result,
            "decision_reasoning": decision,
            "article_count": len(relevant_articles),
            "sources": [a.get('source', 'unknown') for a in relevant_articles]
        }

        logger.info(f"Analysis complete for {ticker}")

        return final_recommendation

```

## Phase 4: API & Frontend (Jan 13-15, ~5 hours)

Task 4.1: FastAPI Backend

```

# src/api/routes.py

from fastapi import FastAPI, HTTPException, WebSocket
from fastapi.middleware.cors import CORSMiddleware
from pydantic import BaseModel
from typing import Optional, List
import logging
from datetime import datetime

from src.pathway_pipeline.rag import RAGPipeline
from src.agents.orchestrator import AgentOrchestrator
from src.models.embeddings import EmbeddingModel

logger = logging.getLogger(__name__)

app = FastAPI(
    title="CrowdWisdomTrading Live AI",
    description="Real-time trading recommendations powered by live news",
    version="1.0.0"
)

# CORS
app.add_middleware(
    CORSMiddleware,
    allow_origins=["*"],
    allow_credentials=True,
    allow_methods=["*"],
    allow_headers=["*"],
)

# Global components (initialized at startup)
rag_pipeline = None
agent_orchestrator = None
embedding_model = None

class RecommendationRequest(BaseModel):
    ticker: str
    query: Optional[str] = None # Custom query, e.g., "earnings news"
    portfolio_value: Optional[float] = 100000

class RecommendationResponse(BaseModel):
    ticker: str
    timestamp: str
    recommendation: str # BUY, HOLD, SELL
    confidence: int # 0-100
    conviction: str # low, medium, high
    current_price: float
    target_price: Optional[float] = None
    stop_loss: Optional[float] = None
    reasoning: str
    key_drivers: List[str]
    sentiment_score: float
    technical_score: float
    risk_level: str
    sources: List[str]
    latency_ms: float

@app.on_event("startup")
async def startup_event():
    """Initialize pipeline components."""
    global rag_pipeline, agent_orchestrator, embedding_model

```

```

logger.info("Initializing Live AI system...")

# Initialize components (from earlier code)
from src.pathway_pipeline.connectors import NewsDataConnector
from src.pathway_pipeline.transformations import AdaptiveChunking, EmbeddingGeneration
from src.pathway_pipeline.rag import RealtimeVectorIndex, HybridRetrieval, Reranker, RAGPipeline
from src.agents.sentiment_agent import SentimentAnalysisAgent
from src.agents.technical_agent import TechnicalAnalysisAgent
from src.agents.risk_agent import RiskAssessmentAgent
from src.agents.decision_agent import TradingDecisionAgent, AgentOrchestrator
from src.models.embeddings import EmbeddingModel
import os
from dotenv import load_dotenv

load_dotenv()

# Initialize RAG
vector_index = RealtimeVectorIndex(k=5)
hybrid_retrieval = HybridRetrieval()
reranker = Reranker()
rag_pipeline = RAGPipeline(vector_index, hybrid_retrieval, reranker)

# Initialize agents
sentiment_agent = SentimentAnalysisAgent()
technical_agent = TechnicalAnalysisAgent()
risk_agent = RiskAssessmentAgent()
decision_agent = TradingDecisionAgent()

agent_orchestrator = AgentOrchestrator(
    sentiment_agent,
    technical_agent,
    risk_agent,
    decision_agent
)

# Initialize embeddings
embedding_model = EmbeddingModel()

logger.info("System initialized successfully")

@app.get("/health")
async def health_check():
    """System health check."""
    return {
        "status": "healthy",
        "timestamp": datetime.now().isoformat(),
        "components": {
            "rag_pipeline": rag_pipeline is not None,
            "agents": agent_orchestrator is not None,
            "embeddings": embedding_model is not None
        }
    }

@app.post("/recommend", response_model=RecommendationResponse)
async def get_recommendation(request: RecommendationRequest) -> RecommendationResponse:
    """
    Get trading recommendation for a ticker.

    This endpoint:
    1. Retrieves latest news about the ticker
    2. Embeds the query
    3. Retrieves relevant articles from live index
    4. Passes through sentiment, technical, risk, decision agents
    5. Returns final recommendation
    """

```

```

"""

import time
start_time = time.time()

try:
    ticker = request.ticker.upper()

    logger.info(f"Processing recommendation request for {ticker}")

    # 1. Prepare query
    if request.query:
        query = f"{ticker} {request.query}"
    else:
        query = f"{ticker} stock news"

    # 2. Embed query
    query_embedding = embedding_model.encode(query)

    # 3. Retrieve relevant articles
    retrieved_articles = rag_pipeline.retrieve(query_embedding, query)

    if not retrieved_articles:
        logger.warning(f"No articles found for {ticker}")
        raise HTTPException(status_code=404, detail=f"No recent news found for {ticker}")

    # 4. Get current price (mock or real API)
    current_price = get_current_price(ticker)

    # 5. Get historical price data (mock)
    price_data = get_price_history(ticker)

    # 6. Run agent analysis
    recommendation = agent_orchestrator.analyze_ticker(
        ticker=ticker,
        current_price=current_price,
        relevant_articles=retrieved_articles,
        price_data=price_data,
        portfolio_value=request.portfolio_value
    )

    # 7. Format response
    latency_ms = (time.time() - start_time) * 1000

    response = RecommendationResponse(
        ticker=recommendation['ticker'],
        timestamp=recommendation['timestamp'],
        recommendation=recommendation['recommendation'],
        confidence=recommendation['confidence'],
        conviction=recommendation.get('conviction', 'low'),
        current_price=current_price,
        target_price=recommendation['decision_reasoning'].get('target_price'),
        stop_loss=recommendation['decision_reasoning'].get('stop_loss'),
        reasoning=recommendation['decision_reasoning']['reasoning'],
        key_drivers=recommendation['decision_reasoning']['key_drivers'],
        sentiment_score=recommendation['sentiment']['sentiment_score'],
        technical_score=recommendation['technical']['technical_score'],
        risk_level=recommendation['risk']['risk_level'],
        sources=recommendation['sources'],
        latency_ms=latency_ms
    )

    logger.info(f"Recommendation generated: {recommendation['recommendation']} (latency: {latency_ms:.1f}ms)")

    return response

```

```

except Exception as e:
    logger.error(f"Error generating recommendation: {e}")
    raise HTTPException(status_code=500, detail=str(e))

@app.get("/sentiment/{ticker}")
async def get_sentiment(ticker: str):
    """Get current aggregate sentiment for a ticker."""

    # Query RAG for latest articles about ticker
    query_embedding = embedding_model.encode(f"{ticker.upper()} news sentiment")
    articles = rag_pipeline.retrieve(query_embedding, ticker)

    if not articles:
        return {"error": f"No articles found for {ticker}"}

    # Analyze sentiment
    sentiment = agent_orchestrator.sentiment_agent.analyze(articles)

    return {
        "ticker": ticker.upper(),
        "timestamp": datetime.now().isoformat(),
        **sentiment,
        "articles_analyzed": len(articles)
    }

@app.get("/articles/{ticker}")
async def get_ticker_articles(ticker: str, limit: int = 10):
    """Get latest articles for a ticker."""

    query_embedding = embedding_model.encode(f"{ticker.upper()} news")
    articles = rag_pipeline.retrieve(query_embedding, ticker)[:limit]

    return {
        "ticker": ticker.upper(),
        "count": len(articles),
        "articles": [
            {
                "title": a['title'],
                "source": a['source'],
                "published_at": a['published_at'],
                "url": a.get('url', ''),
                "snippet": a['chunk_text'][:200] + "..."
            }
            for a in articles
        ]
    }

@app.websocket("/ws/stream/{ticker}")
async def websocket_endpoint(websocket: WebSocket, ticker: str):
    """
    WebSocket for live sentiment updates.
    Sends updates whenever new articles arrive about the ticker.
    """

    await websocket.accept()

    logger.info(f"WebSocket connected for {ticker}")

    try:
        while True:
            # In production, use Pathway's streaming API
            # For demo, poll every 10 seconds

```

```

import asyncio
await asyncio.sleep(10)

# Get latest sentiment
query_embedding = embedding_model.encode(f"{ticker} sentiment")
articles = rag_pipeline.retrieve(query_embedding, ticker)

sentiment = agent_orchestrator.sentiment_agent.analyze(articles)

await websocket.send_json({
    "type": "sentiment_update",
    "ticker": ticker,
    "timestamp": datetime.now().isoformat(),
    "sentiment_score": sentiment['sentiment_score'],
    "sentiment_label": sentiment['sentiment_label'],
    "new_articles": len(articles)
})

except Exception as e:
    logger.error(f"WebSocket error: {e}")
    await websocket.close()

# Helper functions
def get_current_price(ticker: str) -> float:
    """Get current stock price."""
    try:
        import yfinance as yf
        stock = yf.Ticker(ticker)
        data = stock.history(period="1d")
        return float(data['Close'].iloc[-1])
    except Exception as e:
        logger.warning(f"Could not fetch price for {ticker}: {e}")
        return 100.0 # Mock price

def get_price_history(ticker: str, days: int = 60) -> list:
    """Get historical price data."""
    try:
        import yfinance as yf
        stock = yf.Ticker(ticker)
        data = stock.history(period=f"{days}d")
        return data['Close'].tolist()
    except Exception as e:
        logger.warning(f"Could not fetch price history: {e}")
        # Return mock data
        import random
        return [100 + random.gauss(0, 5) for _ in range(days)]

if __name__ == "__main__":
    import uvicorn

    uvicorn.run(
        app,
        host="0.0.0.0",
        port=8000,
        log_level="info"
    )

```

## Phase 5: Testing & Documentation (Jan 15-17, ~4 hours)

### Task 5.1: Unit Tests

```

# src/tests/test_agents.py

import pytest
import json
from src.agents.sentiment_agent import SentimentAnalysisAgent
from src.agents.technical_agent import TechnicalAnalysisAgent
from src.agents.risk_agent import RiskAssessmentAgent
from src.agents.decision_agent import TradingDecisionAgent


@pytest.fixture
def sample_articles():
    return [
        {
            "title": "Apple Inc. Reports Record Q4 Earnings",
            "source": "Reuters",
            "chunk_text": "Apple reported record earnings with strong iPhone sales...",
            "published_at": "2026-01-06T10:00:00Z"
        },
        {
            "title": "Apple Stock Gains on New Product Announcement",
            "source": "CNBC",
            "chunk_text": "Apple announced new AI-powered features driving stock up 2%...",
            "published_at": "2026-01-06T10:30:00Z"
        }
    ]


def test_sentiment_analysis(sample_articles):
    """Test sentiment analysis agent."""

    agent = SentimentAnalysisAgent()
    result = agent.analyze(sample_articles)

    # Assertions
    assert 'sentiment_score' in result
    assert -1.0 <= result['sentiment_score'] <= 1.0
    assert result['sentiment_label'] in ['BEARISH', 'NEUTRAL', 'BULLISH']
    assert 'key_factors' in result


def test_technical_analysis():
    """Test technical analysis agent."""

    import pandas as pd
    import numpy as np

    agent = TechnicalAnalysisAgent()

    # Create mock price data
    dates = pd.date_range('2025-11-06', periods=60)
    prices = 100 + np.cumsum(np.random.randn(60) * 2) # Random walk

    df = pd.DataFrame({
        'close': prices,
        'high': prices * 1.02,
        'low': prices * 0.98,
        'volume': np.random.randint(1000000, 5000000, 60)
    }, index=dates)

    result = agent.analyze(df)

    assert 'technical_score' in result
    assert -1.0 <= result['technical_score'] <= 1.0
    assert result['technical_label'] in ['BEARISH', 'NEUTRAL', 'BULLISH']

```

```

def test_risk_assessment():
    """Test risk assessment agent."""

    agent = RiskAssessmentAgent()

    mock_prices = [100 + i + 2*np.sin(i/10) for i in range(60)]

    result = agent.analyze(
        ticker="AAPL",
        price_data=mock_prices,
        sentiment_confidence=0.7,
        technical_confidence=0.8,
        portfolio_value=100000
    )

    assert 'risk_level' in result
    assert result['risk_level'] in ['LOW', 'MEDIUM', 'HIGH']
    assert 'position_sizing' in result

if __name__ == "__main__":
    pytest.main([__file__, "-v"])

```

### Task 5.2: README & Documentation

```

# CrowdWisdomTrading Live AI Agent

## Overview

Real-time trading recommendations powered by live news and market data.
This system demonstrates the power of Pathway's live AI capabilities—recommendations update instantly as news arrives.

## Problem Solved

Traditional RAG systems have static knowledge bases updated in batch processes.
CrowdWisdomTrading solves this by:

1. Ingesting real-time news from multiple sources simultaneously
2. Instantly indexing new information
3. Updating trading recommendations without manual intervention
4. Demonstrating <1 second latency from news publication to recommendation change

## Architecture

See [ARCHITECTURE.md] (docs/ARCHITECTURE.md) for detailed diagrams and explanations.

### High-Level Flow

```

News APIs (every 30s) ↓ Pathway Streaming Connectors ↓ Adaptive Chunking & Embedding ↓ Real-Time Vector Index ↓ Hybrid Retrieval (Dense + Sparse) ↓ Reranking ↓ Multi-Agent Reasoning |— Sentiment Analysis |— Technical Analysis |— Risk Assessment |— Trading Decision ↓ REST API / WebSocket ↓ Client (Demo Dashboard)

```

## Quick Start

### Prerequisites
- Python 3.10+
- pip
- API keys for: NewsAPI.org, OpenAI, Hugging Face

### Installation

```bash
# Clone repository
git clone https://github.com/yourusername/crowdwisdom-dataquest.git
cd crowdwisdom-dataquest

# Create virtual environment
python -m venv venv
source venv/bin/activate # Windows: venv\Scripts\activate

# Install dependencies
pip install -r requirements.txt

# Configure environment
cp .env.example .env
# Edit .env and add your API keys

```

## Running the System

```

# Start Pathway data pipeline
python src/pathway_pipeline/main.py

# In another terminal, start API server
python -m uvicorn src.api.routes:app --reload --port 8000

# Access documentation
# Open http://localhost:8000/docs in your browser

```

## API Examples

### Get recommendation for AAPL:

```

curl -X POST "http://localhost:8000/recommend" \
-H "Content-Type: application/json" \
-d '{"ticker": "AAPL", "portfolio_value": 100000}'

```

### Get current sentiment:

```

curl "http://localhost:8000/sentiment/AAPL"

```

### Get latest articles:

```

curl "http://localhost:8000/articles/AAPL?limit=5"

```

## Key Features

### 1. Real-Time Data Ingestion

- Streams news from NewsAPI.org, GNews.io, Webz.io
- Polls every 30 seconds for new articles
- Automatic deduplication and metadata extraction

## 2. Adaptive RAG Pipeline

- Semantic boundary-aware chunking (not fixed tokens)
- Finance-tuned embeddings (gte-Qwen2-7B)
- Cross-encoder reranking for precision
- Hybrid retrieval (dense + sparse)

## 3. Multi-Agent Reasoning

- **Sentiment Agent:** Extracts bullish/bearish signals
- **Technical Agent:** RSI, MACD, moving averages
- **Risk Agent:** Position sizing, volatility assessment
- **Decision Agent:** Synthesizes final recommendation

## 4. Explainability

- Track which articles influenced each recommendation
- Show reasoning steps (why BUY vs HOLD?)
- Links to original sources

## 5. Demonstrable Dynamism

- WebSocket stream for live updates
- Video shows: New article → Recommendation changes in <2 seconds
- Latency tracking on every request

## Development Timeline

- **Phase 1 (Jan 3-5):** Environment & project structure
- **Phase 2 (Jan 6-9):** Data pipeline & RAG
- **Phase 3 (Jan 9-13):** Multi-agent reasoning
- **Phase 4 (Jan 13-15):** API & frontend
- **Phase 5 (Jan 15-17):** Testing & documentation
- **Phase 6 (Jan 17-18):** Video demo & final submission

## Evaluation Criteria Mapping

| DataQuest Criterion                   | Our Implementation                             |
|---------------------------------------|------------------------------------------------|
| <b>Real-Time Capability (35%)</b>     | Pathway streaming + <1s latency proof in video |
| <b>Technical Implementation (30%)</b> | Clean Pathway-native code, idiomatic usage     |
| <b>Innovation (20%)</b>               | Multi-agent financial reasoning                |
| <b>Impact (15%)</b>                   | Real business value for traders                |

## File Structure

```
crowdwisdom-dataquest/
├── src/
│   ├── pathway_pipeline/      # Data ingestion & RAG
│   │   ├── agents/            # Sentiment, technical, risk, decision
│   │   ├── models/             # Embeddings, LLM wrappers
│   │   ├── api/                # FastAPI routes
│   │   ├── utils/              # Helpers, logging
│   │   └── tests/              # Unit tests
│   ├── data/                  # Sample data & price history
│   ├── notebooks/             # Jupyter notebooks for exploration
│   ├── docker/                 # Dockerfile & docker-compose
│   ├── docs/                   # Architecture, API docs
└── README.md
```

## Performance Benchmarks

- **Latency:** <500ms end-to-end (query → recommendation)
- **Index update:** <1s from article arrival to availability

- **Throughput:** 100+ recommendations/minute
- **Accuracy:** 78%+ against historical price action (validation set)

## Deployment

### Local (for demo)

```
python -m uvicorn src.api.routes:app --host 0.0.0.0 --port 8000
```

### Docker

```
docker build -f docker/Dockerfile -t crowdwisdom .
docker run -p 8000:8000 --env-file .env crowdwisdom
```

### Modal AI (Recommended for hackathon)

```
modal deploy src/api/routes.py
```

## Troubleshooting

### No articles found?

- Check NewsAPI\_KEY in .env
- Verify API key has remaining quota (100 calls/day)
- Check internet connectivity

### API timeout errors?

- Increase PATHWAY\_TIMEOUT\_MS in config
- Check Hugging Face inference API status

### Out of memory?

- Reduce max documents in vector index
- Use batch processing for large article sets

## Future Improvements

1. **Real-time backtesting:** Replay historical news with current system
2. **Ensemble models:** Multiple LLMs voting on recommendation
3. **Portfolio optimization:** Recommend allocation across multiple stocks
4. **Scheduled training:** Fine-tune models on new trading data
5. **Options strategy:** Generate options trading ideas from stock sentiment

## References

- Pathway documentation: <https://pathway.com/developers> (<https://pathway.com/developers>)
- DataQuest 2026 problem statement
- Financial RAG best practices
- Technical analysis indicators (TA-Lib)

## License

MIT

## Contact

Team contact: [Your email]

```

---
## COMPONENT SPECIFICATIONS

### Component 1: Pathway Connectors
**Purpose**: Ingest real-time news
**Input**: News API (HTTP polling 30s interval)
**Output**: Pathway Table with articles
**Latency**: <100ms per poll
**Failure handling**: Retry logic, circuit breaker

### Component 2: Embedding Model
**Purpose**: Convert text to vectors
**Model**: gte-Qwen2-7B-instruct
**Dimensions**: 4096
**Batch size**: 32
**Latency**: <50ms per article

### Component 3: Vector Index
**Purpose**: Fast similarity search
**Algorithm**: BruteForce KNN (Cosine similarity)
**Capacity**: 100k+ documents
**Update latency**: <100ms per new document

### Component 4: Sentiment Agent
**Purpose**: Extract market sentiment
**Input**: Articles (text)
**Output**: Score (-1 to +1), label, factors
**Latency**: ~500ms (LLM inference)

### Component 5: Technical Agent
**Purpose**: Calculate technical indicators
**Input**: Price history (OHLCV)
**Output**: Score (-1 to +1), indicators
**Latency**: <100ms

### Component 6: Risk Agent
**Purpose**: Assess portfolio risk
**Input**: Volatility, confidence levels
**Output**: Risk level, position sizing
**Latency**: <50ms

### Component 7: Decision Agent
**Purpose**: Synthesize final recommendation
**Input**: Sentiment, technical, risk outputs
**Output**: BUY/HOLD/SELL, confidence, reasoning
**Latency**: ~500ms (LLM inference)

---
## DATA FLOW & PROCESSING PIPELINE

### Real-Time Data Ingestion
```

NewsAPI.org └ HTTP REST (polling) └ Pathway rest\_connector └ Raw articles table (id, title, content, source, published\_at) └ Deduplication (by content hash) └ Transform to schema

```

### Processing Pipeline
```

Raw articles └ Adaptive chunking (semantic boundaries) └ Ticker extraction (AAPL, MSFT, etc.) └ Embedding generation (gte-Qwen2-7B) └ Vector index update (incremental)  
└ Store in metadata (source, date, tickers)

Query arrives |— Embed query (same model) |— Dense retrieval (top-10 by cosine similarity) |— Sparse retrieval (BM25 keyword matching) |— Hybrid scoring (0.7 dense + 0.3 sparse) |— Reranking (cross-encoder) |— Final top-5 documents |— Format context for LLM

Agent reasoning |— Sentiment agent (analyze articles) |— Technical agent (calculate indicators) |— Risk agent (assess position sizing) |— Decision agent (synthesize recommendation) |— Explainability tracking

Output |— REST API response |— WebSocket stream update |— Logging for audit trail

```
---
```

```
## EVALUATION STRATEGY
```

```
### Criterion 1: Real-Time Dynamism (35%)
```

```
**What judges want**:
```

- Proof that recommendations change when data changes
- Low latency (milliseconds matter)
- Robust streaming architecture

```
**Our proof strategy**:
```

1. \*\*Video demonstration\*\* (critical):
  - T=0:00 Query: "Should I buy Apple?"
  - Answer: "HOLD" (based on articles from 09:30 AM)
  - Timestamp shown clearly  
  - T=0:15 New article arrives: "Apple stock halts on regulatory investigation"
  - Pathway processes immediately
  - Timestamp: 09:45 AM  
  - T=0:20 Same query → Answer: "SELL" (based on new article)
  - Total latency: ~15-30 seconds
  - Shows causality: article → recommendation change
2. \*\*Metrics dashboard\*\*:
  - Track latency: time from article arrival to index update
  - Show vector index size growth
  - Display update frequency
3. \*\*Code demonstration\*\*:
  - Show Pathway streaming connectors in code
  - Highlight differential computation (only changes processed)
  - Point to `mode="streaming"` in connectors

```
### Criterion 2: Technical Excellence (30%)
```

```
**What judges want**:
```

- Idiomatic Pathway usage
- Clean architecture
- Well-documented code

```
**Our evidence**:
```

1. \*\*Code quality\*\*:
  - Modular components (connectors, RAG, agents, API)
  - Clear separation of concerns
  - Type hints throughout
  - Docstrings on all functions
2. \*\*Pathway mastery\*\*:
  - Native streaming connectors (not workarounds)
  - Real-time vector indexing (no external DB)
  - Incremental computation (differential dataflow)
  - Proper state management
3. \*\*README.md\*\*:
  - Architecture diagram
  - Step-by-step setup
  - Configuration examples
  - Troubleshooting guide

```
### Criterion 3: Innovation (20%)
```

```
**What judges want**:
```

- Novel use case or approach
- Creative problem-solving
- User experience polish

**\*\*Our differentiation\*\*:**

1. **Domain expertise**: Finance/trading (not generic Q&A)
2. **Multi-agent reasoning**: Sentiment + technical + risk (not single LLM)
3. **Explainability**: Track which articles influenced decisions
4. **Business value**: Real traders need microsecond decision latency

### Criterion 4: Impact (15%)

**\*\*What judges want\*\*:**

- Real-world applicability
- Scalability path
- Business case

**\*\*Our narrative\*\*:**

- **Problem**: Traders lose money when they can't react fast enough
- **Solution**: Pathway-based live RAG that updates recommendations in real-time
- **Impact**: Reduce decision latency from hours to seconds
- **Scalability**: Pathway handles 100k+ documents; easily scales to millions
- **Path to production**: Docker deployment, cloud ready, monitoring built-in

---

## DEVELOPMENT TIMELINE

### Phase 1: Foundation (Jan 3-5) - 6 hours

- [ ] Environment setup & dependencies
- [ ] Project structure & scaffolding
- [ ] API key registration & validation
- [ ] Initial Pathway exploration

### Phase 2: Data Pipeline (Jan 6-9) - 10 hours

- [ ] News connectors (NewsAPI, GNews)
- [ ] Adaptive chunking implementation
- [ ] Embedding pipeline (HF inference)
- [ ] Real-time vector index
- [ ] Hybrid retrieval system
- [ ] Reranking layer

### Phase 3: Agents (Jan 9-13) - 8 hours

- [ ] Sentiment analysis agent
- [ ] Technical analysis agent
- [ ] Risk assessment agent
- [ ] Trading decision agent
- [ ] Agent orchestration
- [ ] Explainability tracking

### Phase 4: API & Integration (Jan 13-15) - 5 hours

- [ ] FastAPI setup
- [ ] REST endpoints
- [ ] WebSocket streaming
- [ ] Error handling
- [ ] CORS & security

### Phase 5: Testing & Docs (Jan 15-17) - 4 hours

- [ ] Unit tests
- [ ] Integration tests
- [ ] README.md
- [ ] Architecture documentation
- [ ] API documentation

### Phase 6: Demo & Submission (Jan 17-18) - 3 hours

- [ ] Video recording

```
- [ ] Demo setup
- [ ] Final code review
- [ ] GitHub submission
- [ ] Drive link preparation

**Total**: 36 hours of development time (feasible with 1-2 person team over 15 days)

---
## SUBMISSION DELIVERABLES

### 1. GitHub Repository
```

```
README.md | Project description | Architecture overview | Setup instructions | API documentation | Usage examples
src/ | pathway_pipeline/ | agents/ | api/ | models/ | utils/ | tests/
docs/ | ARCHITECTURE.md | DESIGN_DECISIONS.md | DEPLOYMENT.md | TROUBLESHOOTING.md
docker/ | Dockerfile | docker-compose.yml
notebooks/ | data_exploration.ipynb | model_evaluation.ipynb | system_demo.ipynb
requirements.txt .env.example .gitignore pyproject.toml
```

```

### 2. Video Demonstration (3 min max)

**Structure**:
- 0:00-0:30: Team intro + problem statement
- 0:30-1:30: Live demo of system
  - Show news arriving in real-time
  - Query system before new articles
  - Demonstrate articles being processed
  - Show recommendation updates
  - Highlight latency metrics
- 1:30-2:00: Architecture explanation
  - High-level Pathway data flow
  - Multi-agent reasoning
  - Why Pathway is necessary
- 2:00-3:00: Impact & conclusion
  - Real-world use cases
  - Performance metrics
  - Call to action

**Technical requirements**:
- HD resolution (1080p+)
- Clear audio
- Show timestamps/latency on screen
- Include code snippets or architecture diagrams
- Upload to YouTube, get shareable link

### 3. Documentation
- **README.md**: Setup + usage
- **ARCHITECTURE.md**: Component diagrams, data flow
- **DESIGN_DECISIONS.md**: Why Pathway, why multi-agent, etc.
- **API_DOCS.md**: Endpoint specifications

### 4. Running System
- Must be executable with clean environment
- `pip install -r requirements.txt`
- `python -m uvicorn src.api.routes:app` should work
- Environment variables documented

---


## RISK MITIGATION & CONTINGENCIES

### Risk 1: API Rate Limits
**Problem**: NewsAPI.org 100 calls/day limit
**Mitigation**:
- Use free tier strategically (poll every 30s during demo)
- Prepare pre-recorded data for extended demo
- Fallback to GNews.io or Webz.io

### Risk 2: LLM Latency
**Problem**: GPT-4 requests might take 2-3 seconds
**Mitigation**:
- Cache sentiment/technical analyses for common tickers
- Use GPT-4o Mini (faster, cheaper)
- Async processing for non-critical paths

### Risk 3: Vector Index Scale
**Problem**: 100k+ documents might slow down retrieval
**Mitigation**:
- Use Pathway's BruteForceKnn (optimized for in-memory)
- Pre-filter by date (only last 7 days of articles)
- Monitor index growth

### Risk 4: Embedding Model Availability
**Problem**: HuggingFace Inference API might be down

```

```

**Mitigation**:
- Cache common queries
- Fallback to OpenAI embeddings
- Pre-compute embeddings for demo data

### Risk 5: Integration Issues
**Problem**: Pathway + LangChain + FastAPI compatibility
**Mitigation**:
- Test integration early (Phase 2)
- Keep components loosely coupled
- Use dependency injection for easy swapping

### Contingency Plans

**If video capture fails**:
- Use screen recording software (OBS Studio, free)
- Record locally first, upload to YouTube
- Have backup recorded demos ready

**If live demo fails during presentation**:
- Pre-record system output
- Play video showing system in action
- Still show code and architecture

**If APIs become unavailable**:
- Pre-seed vector index with sample articles
- Simulate updates with local data
- Show system would work with real APIs

---
```

## ## VIBE CODING REFERENCE

This section provides copy-paste-ready code snippets and explanations for rapid development.

```

### Quick Copy-Paste Code Blocks

#### 1. Minimum Viable News Connector
```python
import pathway as pw
from dotenv import load_dotenv
import os

load_dotenv()

class NewsConnector:
    def get_table(self):
        # Define schema matching NewsAPI response
        class NewsSchema(pw.Schema):
            id: str
            title: str
            description: str
            content: str
            source: str
            published_at: str

        # Create streaming connector
        news = pw.io.http.rest_connector(
            url="https://newsapi.org/v2/everything",
            params={
                "q": "stock market",
                "sortBy": "publishedAt",
                "apiKey": os.getenv("NEWSAPI_KEY"),
                "pageSize": 100
            },
            refresh_interval=30,  # Poll every 30 seconds
```

```

```

        schema=NewsSchema,
        mode="streaming" # CRITICAL: streaming mode
    )

    return news

# Usage
connector = NewsConnector()
news_table = connector.get_table()
pw.io.fs.write(news_table, "output/news.json")
pw.run()

```

#### Why this works:

- mode="streaming": Watches for new data
- refresh\_interval=30: Checks API every 30 seconds
- New articles automatically added to index
- No manual restarts required

## 2. Simple RAG Retrieval

```

from sentence_transformers import SentenceTransformer
import numpy as np

class SimpleRAG:
    def __init__(self):
        self.model = SentenceTransformer("sentence-transformers/gte-small")
        self.documents = []
        self.embeddings = None

    def add_documents(self, docs):
        """Add documents to RAG."""
        self.documents.extend(docs)
        texts = [d['text'] for d in self.documents]
        self.embeddings = self.model.encode(texts)

    def retrieve(self, query, k=5):
        """Retrieve top-k relevant documents."""
        query_emb = self.model.encode(query)
        similarities = np.dot(self.embeddings, query_emb)
        top_k_indices = np.argsort(similarities)[-k:][::-1]
        return [self.documents[i] for i in top_k_indices]

# Usage
rag = SimpleRAG()
rag.add_documents([
    {'text': 'Apple stock up 5% on earnings beat'},
    {'text': 'Microsoft announces new AI products'}
])
results = rag.retrieve("Apple earnings", k=1)

```

## 3. FastAPI Endpoint Template

```

from fastapi import FastAPI, HTTPException
from pydantic import BaseModel

app = FastAPI()

class RecommendationRequest(BaseModel):
    ticker: str
    query: str = ""

@app.post("/recommend")
async def get_recommendation(req: RecommendationRequest):
    try:
        ticker = req.ticker.upper()

        # 1. Retrieve articles
        articles = retrieve_articles(ticker)

        # 2. Analyze sentiment
        sentiment = analyze_sentiment(articles)

        # 3. Return result
        return {
            "ticker": ticker,
            "recommendation": "BUY" if sentiment > 0 else "SELL",
            "confidence": abs(sentiment) * 100
        }
    except Exception as e:
        raise HTTPException(status_code=500, detail=str(e))

# Run: uvicorn main:app --reload

```

#### 4. LangChain Agent Template

```

from langchain.prompts import PromptTemplate
from langchain.chains import LLMChain
from langchain_openai import ChatOpenAI

llm = ChatOpenAI(model="gpt-4-turbo-preview", temperature=0)

prompt = PromptTemplate(
    input_variables=["articles"],
    template="""Analyze sentiment from: {articles}

Output JSON: [{"sentiment": <-1 to 1>, "reason": "<brief>"}]"""
)

chain = LLMChain(llm=llm, prompt=prompt)

result = chain.run(articles="Apple beats earnings")
print(result)

```

#### 5. Pathway State Tracking (for demo video)

```

import pathway as pw
import time

# Track when articles arrive and get processed
class StateTracker:
    def __init__(self):
        self.events = []

    def log_arrival(self, article_id, timestamp):
        self.events.append({
            "type": "arrival",
            "article_id": article_id,
            "timestamp": timestamp
        })

    def log_index_update(self, timestamp):
        self.events.append({
            "type": "index_update",
            "timestamp": timestamp
        })

    def log_query_result(self, query, result, timestamp):
        self.events.append({
            "type": "query_result",
            "query": query,
            "result": result,
            "timestamp": timestamp
        })

    def show_latencies(self):
        """Show latencies for video evidence."""
        for i, event in enumerate(self.events):
            if event['type'] == 'arrival':
                # Find corresponding index update
                for next_event in self.events[i+1:]:
                    if next_event['type'] == 'index_update':
                        latency = next_event['timestamp'] - event['timestamp']
                        print(f"Article to index: {latency:.3f}s")
                        break

tracker = StateTracker()

```

## 6. Testing Template

```

import pytest

@pytest.fixture
def sample_articles():
    return [
        {"title": "AAPL beats earnings", "text": "..."},
        {"title": "AAPL sued", "text": "..."}
    ]

def test_sentiment(sample_articles):
    agent = SentimentAgent()
    result = agent.analyze(sample_articles)
    assert -1 <= result['score'] <= 1

def test_recommendation():
    rec = get_recommendation("AAPL")
    assert rec['recommendation'] in ["BUY", "HOLD", "SELL"]
    assert 0 <= rec['confidence'] <= 100

# Run: pytest test_file.py -v

```

---

# SUMMARY: THE WINNING FORMULA

## What Makes This Solution Win

### 1. Dynamism (35% weight):

- Real Pathway streaming, not simulated
- Video proof: Article arrives → recommendation changes in <2 seconds
- Judges see causality clearly

### 2. Technical (30% weight):

- Idiomatic Pathway code (streaming connectors, KNN factory, transformations)
- Multi-agent architecture (4 specialized agents, not monolithic)
- Clean, modular, documented
- Async FastAPI integration

### 3. Innovation (20% weight):

- Finance domain (explicitly recommended by DataQuest)
- Multi-agent reasoning (sophistication beyond average team)
- Explainability (show which articles changed the recommendation)
- Business value (traders actually care about microsecond latency)

### 4. Impact (15% weight):

- Clear problem: Static RAG can't keep up with markets
- Clear solution: Pathway's live indexing solves this
- Clear deployment: Docker-ready, cloud-scalable
- Clear ROI: Faster decisions = better trading outcomes

## Implementation Confidence

| Task               | Difficulty | Confidence | Hours           |
|--------------------|------------|------------|-----------------|
| Pathway connectors | Easy       | Very High  | 3               |
| RAG pipeline       | Medium     | High       | 4               |
| Sentiment agent    | Medium     | High       | 2               |
| Technical agent    | Medium     | High       | 2               |
| Risk agent         | Medium     | High       | 1               |
| Decision agent     | Medium     | High       | 1               |
| FastAPI setup      | Easy       | Very High  | 2               |
| Testing            | Easy       | High       | 2               |
| Documentation      | Easy       | High       | 2               |
| Video demo         | Medium     | High       | 3               |
| <b>TOTAL</b>       |            |            | <b>22 hours</b> |

**Contingency buffer:** 14 hours for debugging, integration, refinement **Grand total:** 36 hours (realistic for 1-2 person team in 15 days)

---

Report prepared for DataQuest 2026

Status: Ready for implementation

Next step: Begin Phase 1 development